OneRoof: Base-, Variant-, and Consensus-calling under One Proverbial Roof.

Table of contents

`nextflow DSL2` `≥22.10.1` `run with` `docker` `run with` `singularity` `◯ Docker CI` `passing`

Overview

`oneroof` is a Nextflow pipeline designed to take a common series of bioinformatic tasks (see below) and put them under "one roof". We mean this literally as well as figuratively: the pipeline will perform certain tasks best when run on networked devices in the same building.

`oneroof` was originally developed in the early stages of the United States Bovine Highly Pathogenic Avian Influenza (HPAI) outbreak of 2024, when we wanted one, configurable, easy-to-run pipeline that would do all of the following seamlessly:

1. Handle super-accuracy basecalling with GPU acceleration on pod5-formatted Nanopore signal files, working on GCP, AWS, or Slurm if need be.
2. Demultiplex BAM-formatted reads that come out of basecalling.
3. Perform the above two steps as signal files become available, either locally or remotely via an SSH client.
4. Accept raw read BAMs or FASTQs if basecalling and demultiplexing have already been performed elsewhere.
5. Accept paired Illumina reads in addition to Nanopore reads.
6. Use forward and reverse primer sequences to select only those reads that represent complete amplicons.

7. Trim away primers—and also any bases that are upstream of the forward primer or downstream of the reverse primer, *while only retaining those reads identified as complete amplicons by containing both primers.*
8. Align to a custom reference with the proper presets for the provided data.
9. Call variants and consensus sequences with appropriate settings for the provided data.
10. Perform tree building with `nextclade`, quality introspection with `multiQC`, and error correction based on the input sequence platform.

Though many excellent pipelines currently exist, e.g. `nf-core/viralrecon`, `epi2me-labs/wf-amplicon`, and `nf-core/nanoseq`, none of these pipelines quite handled all of the above. `oneroof` seeks to handle these requirements while remaining highly configurable for users, highly modular for developers, and easy to control in the command line for both.

Overall, `oneroof` can be summarized as a variant-calling pipeline written in and managed by Nextflow. Its software dependencies are provided through containers or through an environment assembled by `pixi`. To run it on your own Nanopore pod5s with Docker containers, simply run something like:

```
nextflow run nrminor/oneroof \
--pod5_dir my_pod5_dir \
--primer_bed my_primers.bed \
--refseq my_ref.fasta \
--ref_gbk my_ref.gbk \
--kit "SQK-NBD114-24"
```

These are the core elements required to run on Nanopore data: a directory of pod5 files, a BED file of primer coordinates, a reference sequence in FASTA and Genbank format, and the Nanopore barcoding kit used.

And for Illumina paired-end reads, it's even simpler:

```
nextflow run nrminor/oneroof \
--illumina_fastq_dir my_illumina_reads/
```

If you want to use Apptainer containers instead of Docker, just add `-profile apptainer` to either of the above `nextflow run` commands. And if you don't want to use containers at all, simply run `pixi shell --frozen` to bring all the pipeline's dependencies into scope and then add `-profile containerless` to your `nextflow run` command.

Nextflow pipelines like this one have a ton of configuration, which can be overwhelming for beginners and new users. To make this process easier, we're developing a Terminal User Interface (TUI) to guide you through setup. Please stay tuned!

Quick Start

For most users, `oneroof` will have two core requirements: The Docker container engine, available here, and Nextflow, available here. For users interested in super-accuracy basecalling Nanopore signal files, an on-board GPU supported by the Dorado basecaller is also required.

All remaining software dependencies will be supplied through the pipeline's Docker image, which will be pulled and used to launch containers automatically.

From there, the pipeline's three data dependencies are sequence data in BAM, FASTQ, or POD5 format, a BED file of primer coordinates, and a reference sequence in FASTA and Genbank format. For Nanopore data, a barcoding kit identifier is also required. Simply plug in these files to a command like the above and hit enter!

Detailed Setup Instructions

Configuration

Most users should configure `oneroof` through the command line via the following parameters:

Command line argument	Default value	Explanation
`--primer_bed`	None	A bed file of primer coordinates relative to the reference provided with the parameters `refseq` and `ref_gbk`.
`--refseq`	None	The reference sequence to be used for mapping in FASTA format.
`--ref_gbk`	None	The reference sequence to be used for variant annotation in Genbank format.
`--remote_pod5_location`	None	A remote location to use with an SSH client to watch for pod5 files in realtime as they are generated by the sequencing instrument.
`--file_watcher_config`	None	Configuration file for remote file monitoring. An example can be found at conf/ file_watcher.template.yml.
`--pod5_dir`	None	If a remote pod5 location isn't given, users may provide a local, on-device directory where pod5 files have been manually transferred.
`--precalled_staging`	None	A local directory to watch for Nanopore FASTQs or BAMs as the become available. This is for cases where basecalling is being performed via an-

Command line argument	Default value	Explanation
		other workflow, resulting in BAMs or FASTQs that are gradually transferred into `params.precalled_staging` as basecalling completes.
`--prepped_data`	None	If pod5 files have already been base-called, demultiplexed, and transferred to a local directory accessible to onerood users can specify their location with `--prepped_data`.
`--illumina_fastq_dir`	None	If users have Illumina data to be processed, they may specify their paired-end FASTQ files' location with `illumina_fastq_dir`.
`--model`	sup@latest	the Nanopore basecalling model to apply to the provided pod5 data (defaults to the latest super-accuracy version)
`--kit`	None	The Nanopore barcoding kit used to prepare sequencing libraries.
`--pod5_batch_size`	all pod5s	How many pod5 files to basecall at once. With a single available GPU, all pod5 files should be basecalled together, so this parameter defaults to telling Nextflow to take all pod5 files at once.
`--basecall_max`	1	If basecalling pod5 files is to be parallelized across multiple available GPUs, this parameter tells Nextflow how many parallel instances of the basecaller to run at once (defaults to 1).
`--max_len`	12345678910	The maximum acceptable length for a given read.
`--min_len`	1	The minimum acceptable length for a given read.
`--min_qual`	0	The minimum acceptable average quality for a given read.
`--max_mismatch`	0	The maximum number of mismatches to allow when finding primers

Command line argument	Default value	Explanation
`--downsample_to`	0	Desired coverage to downsample to, with a special value of 0 to indicate 0 downsampling
`--secondary`	None	Whether to turn on secondary alignments for each amplicon. Secondary alignments can increase depth at the cost of more reads potentially mapping to the wrong locations. By default, secondary alignments are off.
`--min_consensus_freq`	0.5	The minimum required frequency of a variant base to be included in a consensu sequence.
`--min_depth_coverage`	20	Minimum required depth of coverage to call a variant.
`--min_variant_frequency`	0.05 (illumina) or 0.10 (nanopore)	Minimum variant frequency to call a variant.
`--meta_ref`	None	Dataset, either a local FASTA file or a pre-built dataset built by Sylph, to use for metagenomic profiling. Can download pre-built ones here: Pre-built Sylph Databases.
`--sylph_tax_db`	None	The taxonomic annotation for the sylph database specified with `--meta_ref`. The pipeline automaticially downloads the databases so only the identifier is needed here.
`--nextclade_dataset`	None	The name of the dataset to run nextclade with. To see all dataset options run `nextclade dataset list --only-names`.
`--results`	results/	Where to place results.
`--cleanup`	false	Whether to cleanup work directory after a successful run.

Note that `oneroof` checks for how to gather data in a particular order for Nanopore data, which mirrors the order they are listed in the table above:

1. First, `oneroof` checks if the user has supplied a remote location with `--remote_pod5_location`, at which point it will launch the file watcher module and begin transferring and basecalling

batches of pod5 files. The size of these batches can be controlled with `--pod5_batch_size`, which defaults to basecalling batches of 100 pod5s.

2. If no remote directory is provided, `oneroof` checks for a local pod5 directory from `--pod5_dir`.
3. If the user doesn't provide a local `--pod5_dir`, `oneroof` assumes that pre-basecalled BAMs or FASTQs are being provided. These can either be watched for with the directory from `--precalled_staging`, or run immediately with the directory from `--prepped_data`.

Developer Setup

`oneroof` depends on software packages supplied through various `conda` registries as well as through `PyPI`, the Python Package Index. To unify these various channels, we used the relatively new `pixi` package and environment manager, which stores all dependencies from both locations in the file `pyproject.toml`.

To reproduce the environment required by this pipeline, make sure you are on a Mac, a linux machine, or a Windows machine using Windows Subsystem for Linux. Then, to reproduce the environment, install pixi with:

```
PIXI_ARCH=x86_64 curl -fsSL https://pixi.sh/install.sh | bash
```

Download the pipeline with:

```
git clone https://github.com/nrminor/oneroof.git && cd oneroof
```

And then open a `pixi` subshell within your terminal with:

```
pixi shell --frozen
```

As long as you are using a supported system, the pipeline should run within that subshell. You can also run the pipeline within that subshell without containers using the "containerless" profile:

```
nextflow run . \
-profile containerless \
--pod5_dir my_pod5_dir \
--primer_bed my_primers.bed \
--refseq my_ref.fasta \
--ref_gbk my_ref.gbk \
--kit "SQK-NBD114-24"
```

Especially on Apple Silicon Macs, this will reduce the overhead of using the Docker Virtual Machine and allow the pipeline to invoke tools installed directly within the local project environment.

Note also that more information on the repo's files is available in our developer guide.

Testing

OneRoof includes a comprehensive test suite built with nf-test, the official testing framework for Nextflow pipelines. The test suite validates pipeline functionality through module, workflow, and end-to-end tests.

Running Tests

Before contributing or deploying changes, ensure all tests pass:

```
# Ensure you're in the Pixi environment (includes nf-test)
pixi shell --frozen

# Run all tests
just test

# Run tests with verbose output
just test-verbose

# Run specific test categories
just test-modules      # Test individual processes
just test-workflows    # Test platform workflows
just test-pipeline     # Test end-to-end functionality

# Run a specific test file
just test-file tests/modules/minimap2_align.nf.test

# Update test snapshots after intentional changes
just test-update
```

Test Structure

Tests are organized in the `tests/` directory: - `tests/modules/` - Unit tests for individual processes - `tests/workflows/` - Integration tests for platform-specific workflows - `tests/pipelines/` - End-to-end pipeline tests - `tests/data/` - Minimal test datasets

The test suite uses minimal synthetic data to ensure fast execution while still exercising all key pipeline features. Tests are automatically run in CI/CD on all pull requests and pushes to main branches.

For more details on the testing framework and how to write new tests, see the test suite documentation.

Further Documentation

For more detailed information about specific aspects of the OneRoof pipeline, please refer to the following documentation:

- **Developer Guide** - Comprehensive guide for developers working with the codebase, including project structure, coding standards, and development workflows

- **Pipeline Architecture** - Detailed technical documentation of the pipeline's architecture, including workflow diagrams, module descriptions, and data flow
- **File Reference Guide** - Complete listing of all files in the repository with descriptions of their purpose and functionality
- **Globus Integration** - Instructions for setting up and using Globus integration for automated data transfer and remote workflow execution
- **Test Suite Documentation** - Guide to the testing framework, including how to write and run tests, test data organization, and CI/CD integration

Contributing

Contributions, feature requests, improvement suggestions, and bug reports via GitHub issues are all welcome! For more information on how to work with the project and what all the repo files are, see our developer guide.

Is it any good?

Yes.

Citation

Lail, Andrew J., William C. Vuyk, Heather Machkovech, Nicholas R. Minor, Nura R. Hassan, Rhea Dalvie, Isla E. Emmen et al. "Amplicon sequencing of pasteurized retail dairy enables genomic surveillance of H5N1 avian influenza virus in United States cattle." PloS one 20, no. 6 (2025): https://doi.org/10.1371/journal.pone.0325203.

```
@article{Lail2025-xf,
  title = "Amplicon sequencing of pasteurized retail dairy enables genomic
surveillance of {H5N1} avian influenza virus in United States cattle",
  author = "Lail, Andrew J and Vuyk, William C and Machkovech, Heather and
Minor, Nicholas R and Hassan, Nura R and Dalvie, Rhea and Emmen, Isla E and
Wolf, Sydney and Kalweit, Annabelle and Wilson, Nancy and Newman, Christina M
and Tiburcio, Patrick Barros and Weiler, Andrea and Friedrich, Thomas C and
O'Connor, David H",
  journal = "PLoS One",
  publisher = "Public Library of Science (PLoS)",
  volume =  20,
  number =  6,
  pages = "e0325203",
  month =  jun,
  year =  2025,
  copyright = "http://creativecommons.org/licenses/by/4.0/",
  language = "en"
}
```

Kwon, Taeyong, Jessie D. Trujillo, Mariano Carossino, Heather M. Machkovech, Konner Cool, Eu Lim Lyoo, Gagandeep Singh et al. "Pathogenicity and transmissibility of bovine-derived HPAI H5N1 B3. 13 virus in pigs." Emerging Microbes & Infections just-accepted (2025): https://doi.org/10.1080/22221751.2025.2509742.

```
@ARTICLE{Kwon2025-yq,
  title = "Pathogenicity and transmissibility of bovine-derived {HPAI} {H5N1}
B3.13 virus in pigs",
  author = "Kwon, Taeyong and Trujillo, Jessie D and Carossino, Mariano and
Machkovech, Heather M and Cool, Konner and Lyoo, Eu Lim and Singh, Gagandeep
and Kafle, Sujan and Elango, Shanmugasundaram and Vediyappan, Govindsamy and
Wei, Wanting and Minor, Nicholas and Matias-Ferreyra, Franco S and Morozov,
Igor and Gaudreault, Natasha N and Balasuriya, Udeni B R and Hensley, Lisa E
and Diel, Diego G and Ma, Wenjun and Friedrich, Thomas C and Richt, Juergen
A",
  journal = "Emerg. Microbes Infect.",
  publisher = "Informa UK Limited",
  volume =  14,
  number =  1,
  pages = "2509742",
  month =  dec,
  year =  2025,
  keywords = "Cattle; H5N1 genotype B3.13; highly pathogenic avian influenza;
mammalian-like mutation; pathogenicity; pig; transmissibility",
  copyright = "http://creativecommons.org/licenses/by-nc/4.0/",
  language = "en"
}
```